

Smart-Cache: Pre-fetching Using Markov Chains in Distributed Caching

Meia Alsup*

Massachusetts Institute of Technology
Cambridge, MA

Amir Farhat*

Massachusetts Institute of Technology
Cambridge, MA

Alexander J. Root*

Massachusetts Institute of Technology
Cambridge, MA

ABSTRACT

Current caching methods for distributed data-stores do not detect patterns in data accesses. Many types of data accesses, such as web page accesses and machine learning workloads have simple patterns that can be detected via a Hidden Markov Model (HMM). This work presents Smart-Cache, a distributed caching model that uses HMMs to predict future data accesses and prefetch data according to some learned patterns. We have implemented this cache and compared it to a memcached-like cache that does no data prefetching. Smart-Cache provides significant speed up on workloads with patterns. The experiments conducted show execution time reduction of 86% for deterministic workload

KEYWORDS

distributed systems, caching, markov chain, distributed learning

1 INTRODUCTION AND RELATED WORK

Efficient, distributed caching is a well-known solution to alleviating load on underlying and expensive-to-query data stores [8]. Some caching systems attempt to predict future data accesses and pre-fetch accordingly. Previous efforts to optimize pre-fetch mechanisms have relied on assumptions around data locality [1]; few methods relax the locality assumption for truly un-opinionated cache pre-fetch prediction.

Memcached also uses Consistent Hashing to distribute load across caches [5]. Consistent Hashing is complex and makes Memcached resilient to hot files, and dynamically reallocates responsibility for files amongst caches in a cache-locality aware way. Smart-Cache takes a much simpler approach to distribution of file responsibility amongst caches, which greatly reduces the overhead while still providing good performance.

Markov chains have been shown to learn data access patterns in many other settings including file I/O, hard disks access, and in smaller scale buffers or pages [4, 6, 7]. Smart-cache is motivated in particular by the deep learning workload scenario, but is generalizable and usable by a much wider array of scenarios with patterned workload accesses.

Machine learning workloads, which are traditionally bottlenecked by compute power and GPU cycles, benefit from new hardware so as to put increasing load on underlying datastores. With state-of-the-art hardware, these datastores become a significant bottleneck [9]. In such situations, even a simple caching layer in front of the datastore gives huge reductions in datastore egress requirements

and significantly lowers round-trip request latency [9]. When training large deep neural networks, most hyper parameter optimization procedures require that a large set of training runs are started, each with different hyper-parameter configurations. While each run has a different configuration, all runs access data in the same order. These runs, which are often in the same cluster of GPUs, all make queries to data with the same access pattern.

Our contributions are three fold. First, we introduce Smart-Cache, a fault-tolerant and distributed cache based on memcached [3] that can dynamically learn data access patterns. Second, we provide a simple patterned workload generator to proxy web workloads. Third, we characterize the performance of Smart-Cache on three workloads with varying degrees of contained patterns, and demonstrate significant improvement with respect to both latency measurements and queries to the underlying datastore.

2 SYSTEM AND SETUP

2.1 Design

2.1.1 Assumptions. We assume a static underlying datastore. Smart-Cache takes initial inspiration from machine learning workloads. ML workloads are short-lived on a static underlying dataset, where modifications to the underlying datastore only occur between task runs. In the first iteration of Smart-Cache, a static underlying datastore is assumed, which enables a simple design.

We also assume similar file popularity amongst files in the datastore. This version of Smart-Cache is not resilient to hot files and nodes.

Extensions to Smart-Cache to mitigate the effects of both of these assumptions are discussed in Section 4.

2.1.2 System Design. Figure 1 shows the system components. Smart-Cache supports a single task at a time, which coordinates many clients. The task interacts with *Cache master* which coordinates the behavior of the individual caching nodes. *Caching nodes* cache and store data, and are the only nodes connected to the underlying datastore.

The number of cache nodes, k , and the replication factor r , are configurable. These two parameters determine the number of file groups, $G = \frac{k}{r}$. In the event r is not divisible by k , the floor is used and extra caches are allocated to file groups at random. These parameters allow the user to scale the caches along two axes. In Figure 1, k is 15, r is 3, and the number of file groups G is 5. File groups are color coded. The cache master (CM) handles distributing files across the G file groups. Consider the underlying data-set \mathcal{D} and $\{d_1, \dots, d_g\}$ as disjoint subsets of \mathcal{D} such that $\mathcal{D} = \bigcup_{i=1}^n d_i$. For each subset of $d_i \subseteq \mathcal{D}$, a file group is responsible for that subset and the subset is replicated across r caching nodes.

*Denotes equal contribution.

2.3 Markov Prefetching Model

Each cache stores a model of file access patterns for the particular cache. We use a sparse Markov model [10], as it is expected that a series of accesses with commonly occurring patterns is best represented in a sparse model. Assuming constant degree d_i for each of the n nodes, $i \in [0, n)$, then a sparse representation takes $\Theta(n)$ memory as opposed to a dense model that requires $\Theta(n^2)$ memory.

Using this model, we can implement an efficient predictor that takes a source file and can predict the next k file accesses in $\Theta(k \log(k))$ time, using a modification of Dijkstra’s algorithm [2]. The problem formulation involves a reduction from an attempt to maximize the product of probabilities in the range $(0, 1]$ to a maximization of the logarithm of the product, which reduces to a minimization of the sum of the logarithm of the probabilities, which reduces to a formulation of Dijkstra’s, as the edge weights are strictly non-negative. It is important to not that this algorithm is aggressively sublinear, as the predictor does not depend on the number of files n that a given cache is responsible for.

3 RESULTS

The results from our runs in time are shown in Table 1. Hits and misses to the underlying datastore for each of our runs are shown in Table 2.

A limitation of the results is that these workloads were run locally on our machines, instead of on a dedicated compute cluster. These results are also for a single run rather than an average across several runs. As such, the difference observed in the Random task between LRU and Markov Chain is likely not statistically significant. However, the result on the Deterministic task is likely significant.

Table 1: Total time in seconds to perform tasks.

	LRU	Markov Chain
Random	11.11	19.12
Patterned	30.01	27.46
Deterministic	26.53	20.47

Table 2: Hits, Misses, and Datastore Queries

Workload Type	LRU		Markov Chain	
	Hits	Misses	Hits	Misses
Random	404	2596	400	2600
Patterned	1112	1062	1638	536
Deterministic	9000	1000	9898	102

4 CONCLUSION AND FUTURE WORK

In this paper, we demonstrate the utility of Smart-Cache to speed up workloads that exhibit patterns. This is immediately useful for machine learning applications and in any other application where workloads can be characterized by repeat data access patterns.

This can be expanded upon on a few axes. The cache master can dynamically allocate cache nodes across file groups to prevent

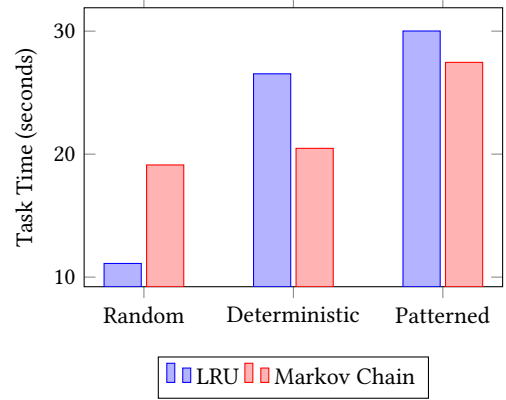


Figure 2: Total time to perform tasks, in seconds, across workloads and caching strategies.




Figure 3: % speedup to perform tasks between caching strategies for each workload

hot files from taking the system down and make the system more resilient to caches going down in a particular file group. The system currently uses a simple hashing scheme, but more complex hashing schemes such as consistent hashing [5] can be leveraged to make the system more tolerant to caches going off and online.

Further, the cache master could implement periodic Markov chain syncing across caches in the same file group. This would enable learning to happen quicker, since patterns observed can be learned at double or triple the speed for a replication factor of two or three respectively.

Another extension to Smart-Cache is supporting a dynamic underlying datastore. This requires the cache-master to allocate and de-allocate files as they are added and deleted from the datastore. With respect to staleness, a system leveraging version numbers could be used, or writes could be set up to notify all caches in the affected file group that the data is stale and in need of refresh.

ACKNOWLEDGMENTS

We’d like to thank Professor Robert Morris, Anish Athalye, and the entire 6.824 staff for guidance and helpful conversations. We would also like to thank Jacob Kahn for providing insight into machine

learning workloads at Facebook AI Research, and inspiring this work.

A APPENDIX

A.1 Code

Our released codebase for Smart-Cache as well as task and workload generation can be found on our Github: <https://github.com/rootjalex/smart-cache>.

A.2 Patterned Workloads

Patterned workloads are generated by first creating short access patterns that will be common amongst clients. In our test runs, we generate 10-200 patterns, depending on the benchmark. From this set of patterns, the workload for each client is created by randomly sampling a pattern from the pattern set 2-10 times, and concatenating them.

To make this concrete, the following example illustrates a potential outcome with two patterns and two client workloads each composed of five randomly concatenated patterns.

- The two randomly generated patterns are [A.txt, B.txt, C.txt] and [X.txt, Y.txt].
- Client 1 has workload: [A.txt, B.txt, C.txt, A.txt, B.txt, C.txt, X.txt, Y.txt, X.txt, Y.txt]
- Client 2 has workload: [X.txt, Y.txt, A.txt, B.txt, C.txt, X.txt, Y.txt, X.txt, Y.txt, A.txt, B.txt, C.txt]

A.3 Cache Responsibility - Hashing Function

The hash function generator takes the following inputs: number of caches (k), cache replication factor (r), list of client IDs, and filenames (n files total) in the underlying datastore. The following steps generate the hash function:

- (1) The number of file groups, G , is determined from $\frac{k}{r}$. Caches are allocated evenly across file groups such that each group is responsible for $[r, r + 1]$ caches in the event of overflow.
- (2) Responsibility for the n files is divided across the G file groups evenly such that each file group is responsible for $[\lfloor \frac{n}{G} \rfloor, \lfloor \frac{n}{G} \rfloor + 1]$ files. Responsibility is split randomly leveraging Go's Random package¹.
- (3) A mapping is created from (client ID, filegroup) \rightarrow [list of cache IDs] for all pairs of client IDs and file names. For every file group and for every client ID, the list of cache IDs assigned to that file group is shuffled randomly. The ordering determined here is the order the client will attempt to retrieve files from caches. This random shuffling helps ensure that no one cache in any file group is fielding too many requests at once. A round robin method was also considered. However, with round robin, if one cache goes down, the cache immediately downstream becomes responsible for twice as many clients. Shuffling mitigates this problem since clients assigned to a dead cache redistribute amongst the remaining alive caches at random. While this is not necessarily

exactly an even split, in a large system with many caches and clients, the split is expected to be almost even and sufficient.

REFERENCES

- [1] Swapnil Bhatia, Elizabeth Varki, and Arif Merchant. 2010. Sequential Prefetch Cache Sizing for Maximal Hit Rate. 89–98. <https://doi.org/10.1109/MASCOTS.2010.18>
- [2] Edsger W Dijkstra. 1959. A note on two problems in connexion with graphs. *Numerische mathematik* 1, 1 (1959), 269–271.
- [3] Brad Fitzpatrick. 2004. Distributed caching with memcached. *Linux journal* 2004, 124 (2004), 5.
- [4] Doug Joseph and Dirk Grunwald. 1997. Prefetching Using Markov Predictors. In *Proceedings of the 24th Annual International Symposium on Computer Architecture* (Denver, Colorado, USA) (ISCA '97). Association for Computing Machinery, New York, NY, USA, 252–263. <https://doi.org/10.1145/264107.264207>
- [5] David Karger, Eric Lehman, Tom Leighton, Rina Panigrahy, Matthew Levine, and Daniel Lewin. 1997. Consistent Hashing and Random Trees: Distributed Caching Protocols for Relieving Hot Spots on the World Wide Web. In *Proceedings of the Twenty-Ninth Annual ACM Symposium on Theory of Computing* (El Paso, Texas, USA) (STOC '97). Association for Computing Machinery, New York, NY, USA, 654–663. <https://doi.org/10.1145/258533.258660>
- [6] A. Laga, J. Boukhobza, M. Koskas, and F. Singhoff. 2016. Lynx: a learning linux prefetching mechanism for SSD performance model. , 6 pages.
- [7] Mingju Li, Elizabeth Varki, Swapnil Bhatia, and Arif Merchant. 2008. TaP: Table-Based Prefetching for Storage Caches. In *Proceedings of the 6th USENIX Conference on File and Storage Technologies* (San Jose, California) (FAST'08). USENIX Association, USA, Article 6, 16 pages.
- [8] Rajesh Nishtala, Hans Fugal, Steven Grimm, Marc Kwiatkowski, Herman Lee, Harry C. Li, Ryan McElroy, Mike Paleczny, Daniel Peek, Paul Saab, David Stafford, Tony Tung, and Venkateshwaran Venkataramani. 2013. Scaling Memcache at Facebook. In *Proceedings of the 10th USENIX Conference on Networked Systems Design and Implementation* (Lombard, IL) (nsdi'13). USENIX Association, USA, 385–398.
- [9] Vineel Pratap Vitaliy Liptchinsky, Jacob Kahn. 2020. Scaling Deep Learning for Automatic Speech Recognition. (2020). <https://developer.nvidia.com/gtc/2020/video/s21838> S21838.
- [10] Jie Xiong, Väinö Jääskinen, and Jukka Corander. 2016. Recursive Learning for Sparse Markov Models. *Bayesian Anal.* 11, 1 (03 2016), 247–263. <https://doi.org/10.1214/15-BA949>

¹<https://golang.org/pkg/math/rand/>